

Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: grader_weights(), grader_sigmoid(), grader_logloss() etc, you should not change those function definition.

Every Grader function has to return True.

Importing packages

```
In [1]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn import linear_model
import math
import random
import matplotlib.pyplot as plt
```

```
In [2]: import numpy as np
```

Creating custom dataset

```
In [3]: # please don't change random_state
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_re
                        n_classes=2, weights=[0.7], class_sep=0.7, random_stat
# make_classification is used to create custom dataset
# Please check this link (https://scikit-learn.org/stable/modules/generated/skLea
```

```
In [4]: X.shape, y.shape
```

```
Out[4]: ((50000, 15), (50000,))
```

Splitting data into train and test

```
In [5]: #please don't change random state
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_
```

```
In [6]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
Out[6]: ((37500, 15), (37500,)), (12500, 15), (12500,))
```

SGD classifier

```
In [7]: # alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial Learning rate for the 'constant', 'invscaling' or 'adaptive' schedu

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_st
clf
# Please check this documentation (https://scikit-learn.org/stable/modules/genera
```

```
Out[7]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0001,
    fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
    loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
    penalty='l2', power_t=0.5, random_state=15, shuffle=True,
    tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
In [8]: clf.fit(X=X_train, y=y_train) # fitting our model
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.02 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.03 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.04 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.05 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.06 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.07 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.08 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.09 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.10 seconds.
Convergence after 10 epochs took 0.10 seconds
```

```
Out[8]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                    early_stopping=False, epsilon=0.1, eta0=0.0001,
                    fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                    loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                    penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                    tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
In [9]: clf.coef_, clf.coef_.shape, clf.intercept_
#clf.coef_ will return the weights
#clf.coef_.shape will return the shape of weights
#clf.intercept_ will return the intercept term
```

```
Out[9]: (array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
                  0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
                  0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
         (1, 15),
         array([ -0.8531383]))
```

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

1. We will be giving you some functions, please write code in that functions only.
2. After every function, we will be giving you expected output, please make sure that you get that output.

- Initialize the weight_vector and intercept term to zeros (Write your code in `def initialize_weights()`)
- Create a loss function (Write your code in `def logloss()`)

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

- for each epoch:
 - for each batch of data points in train: (keep batch size=1)
 - calculate the gradient of loss function w.r.t each weight in weight vector (write your code in `def gradient_dw()`)

$$dw^{(t)} = x_n(y_n - \sigma((w^{(t)})^T x_n + b^t)) - \frac{\lambda}{N} w^{(t)}$$
 - Calculate the gradient of the intercept (write your code in `def gradient_db()`) [check this \(https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing\)](https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing)

$$db^{(t)} = y_n - \sigma((w^{(t)})^T x_n + b^t)$$
 - Update weights and intercept (check the equation number 32 in the above mentioned [pdf \(https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing\)](https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing)):

$$w^{(t+1)} \leftarrow w^{(t)} + \alpha(dw^{(t)})$$

$$b^{(t+1)} \leftarrow b^{(t)} + \alpha(db^{(t)})$$
 - calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
 - And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
 - append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)

Initialize weights

```
In [10]: def initialize_weights(dim):
    ''' In this function, we will initialize our weights and bias'''
    #initialize the weights to zeros array of (dim,1) dimensions
    #you use zeros_like function to initialize zero, check this link https://docs
    #initialize bias to zero
    w = np.zeros(len(dim))
    b = 0
    return w,b
```

```
In [11]: dim=X_train[0]
w,b = initialize_weights(dim)
print('w =',(w))
print('b =',str(b))
```

```
w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Grader function - 1

```
In [12]: dim=X_train[0]
w,b = initialize_weights(dim)

#assert - Test if a condition returns True, if not, the program will raise an Ass
def grader_weights(w,b):
    assert((len(w)==len(dim)) and b==0 and np.sum(w)==0.0)
    return True
grader_weights(w,b)
```

Out[12]: True

Compute sigmoid

$$\text{sigmoid}(z) = 1/(1 + \exp(-z))$$

```
In [13]: def sigmoid(z):
    ''' In this function, we will return sigmoid of z'''
    # compute sigmoid(z) and return
    sigmoid_value=1/(1+np.exp(-z))

    return sigmoid_value
```

Grader function - 2

```
In [14]: def grader_sigmoid(z):
    val=sigmoid(z)
    assert(val==0.8807970779778823)
    return True
grader_sigmoid(2)
```

Out[14]: True

Compute loss

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

```
In [15]: def logloss(y_true, y_pred):

    loss = 0
    for index in range(len(y_true)):
        a = (y_true[index] * math.log(y_pred[index], 10)) + (1 - y_true[index]) *
        b = (-1/len(y_true))
        loss = loss + a * b
    return loss
```

Grader function - 3

```
In [16]: def grader_logloss(true,pred):
    loss=logloss(true,pred)
    #print(loss)
    assert(loss==0.07644900402910387)
    return True
true=[1,1,0,1,0]
pred=[0.9,0.8,0.1,0.8,0.2]
grader_logloss(true,pred)
```

Out[16]: True

Compute gradient w.r.to 'w'

$$dw^{(t)} = x_n(y_n - \sigma((w^{(t)})^T x_n + b^t)) - \frac{\lambda}{N} w^{(t)}$$

```
In [17]: def gradient_dw(x,y,w,b,alpha,N):
    '''In this function, we will compute the gardient w.r.to w '''
    error = y - sigmoid(np.dot(x, w.T) + b)
    dw = x * error
    return dw
```

Grader function - 4

```
In [18]: def grader_dw(x,y,w,b,alpha,N):
    grad_dw=gradient_dw(x,y,w,b,alpha,N)
    assert(np.sum(grad_dw)==2.613689585)
    return True
grad_x=np.array([-2.07864835, 3.31604252, -0.79104357, -3.87045546, -1.14783286,
                -2.81434437, -0.86771071, -0.04073287, 0.84827878, 1.99451725,
                3.67152472, 0.01451875, 2.01062888, 0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)
```

Out[18]: True

Compute gradient w.r.to 'b'

$$db^{(t)} = y_n - \sigma((w^{(t)})^T x_n + b^t))$$

```
In [19]: def gradient_db(x,y,w,b):

    '''In this function, we will compute gradient w.r.to b '''
    db = y - sigmoid(np.dot(x, w.T) + b)
    return db
```

Grader function - 5

```
In [20]: def grader_db(x,y,w,b):
    grad_db=gradient_db(x,y,w,b)
    assert(grad_db==0.5)
    return True

grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,
                -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
                3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_db(grad_x,grad_y,grad_w,grad_b)
```

Out[20]: True

```
In [21]: def probability(x, w, b):
    predicted = []
    probability = []
    for i in range(len(x)):
        z = np.dot(w, x[i]) + b
        sig = sigmoid(z)
        if sig >= 0.5:
            predicted.append(1)
        else:
            predicted.append(0)
        probability.append(sig)
    return np.array(predicted), np.array(probability)

def find_accuracy(actual, predicted):
    return round((np.sum((actual == predicted) /
                        len(actual))) * 100, 2)
```

In []:

In []:

Implementing logistic regression

```

In [22]: def train(X_train,y_train,X_test,y_test,epochs,alpha,eta0):
''' In this function, we will implement logistic regression'''
#Here eta0 is learning rate
#implement the code as follows
# initialize the weights (call the initialize_weights(X_train[0]) function)
# for every epoch
    # for every data point(X_train,y_train)
        #compute gradient w.r.to w (call the gradient_dw() function)
        #compute gradient w.r.to b (call the gradient_db() function)
        #update w, b
    # predict the output of x_train[for all data points in X_train] using w,b
    #compute the loss between predicted and actual values (call the loss func
    # store all the train loss values in a list
    # predict the output of x_test[for all data points in X_test] using w,b
    #compute the loss between predicted and actual values (call the loss func
    # store all the test loss values in a list
    # you can also compare previous loss and current loss, if loss is not upc

global train_loss
train_loss = []

global test_loss
test_loss = []

w, b = initialize_weights(X_train[0])
for ep in range(epochs):
    print('Epoch:', ep + 1)
    for index in range(len(X_train)):
        r_index = random.randint(0, len(X_train) - 1)

        ln_eqn = np.dot(X_train[r_index], w.T) + b
        error = y_train[r_index] - sigmoid(ln_eqn)

        dw = X_train[r_index] * error
        db = error

        w = w + (alpha * dw)
        b = b + (alpha * db)

    predicted, score = probability(X_train, w, b)
    loss = logloss(y_train, score)
    train_loss.append(loss)
    tr_acc = find_accuracy(y_train, predicted)

    predicted, score = probability(X_test, w, b)
    loss = logloss(y_test, score)
    test_loss.append(loss)
    te_acc = find_accuracy(y_test, predicted)

    print('Train Loss:', train_loss[ep], ', Test Loss:', test_loss[ep])
    print('Train Accuracy:', tr_acc, ', Test Accuracy:', te_acc)

return w, b

```



```
In [23]: alpha=0.0001
eta0=0.0001
N=len(X_train)
epochs=10
w,b=train(X_train,y_train,X_test,y_test,epochs,alpha,eta0)
```

```
Epoch: 1
Train Loss: 0.17556467020664712 , Test Loss: 0.17589712592537024
Train Accuracy: 82.55 , Test Accuracy: 82.73
Epoch: 2
Train Loss: 0.16871589309416066 , Test Loss: 0.16931669506625202
Train Accuracy: 82.78 , Test Accuracy: 82.92
Epoch: 3
Train Loss: 0.1663714997367262 , Test Loss: 0.16714362384052414
Train Accuracy: 83.04 , Test Accuracy: 83.1
Epoch: 4
Train Loss: 0.16552499133878631 , Test Loss: 0.16637414926523583
Train Accuracy: 83.13 , Test Accuracy: 83.13
Epoch: 5
Train Loss: 0.1649454756130153 , Test Loss: 0.16586918423537553
Train Accuracy: 83.19 , Test Accuracy: 83.18
Epoch: 6
Train Loss: 0.1646199489560374 , Test Loss: 0.16546320847261684
Train Accuracy: 83.16 , Test Accuracy: 83.2
Epoch: 7
Train Loss: 0.16454998901589163 , Test Loss: 0.1655462546647139
Train Accuracy: 83.15 , Test Accuracy: 83.35
Epoch: 8
Train Loss: 0.16443675893174986 , Test Loss: 0.16536728321888386
Train Accuracy: 83.12 , Test Accuracy: 83.3
Epoch: 9
Train Loss: 0.16447608254263127 , Test Loss: 0.16549857699013598
Train Accuracy: 83.19 , Test Accuracy: 83.27
Epoch: 10
Train Loss: 0.16461465920619656 , Test Loss: 0.16531085658813208
Train Accuracy: 83.11 , Test Accuracy: 83.22
```

Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of 10^{-3}

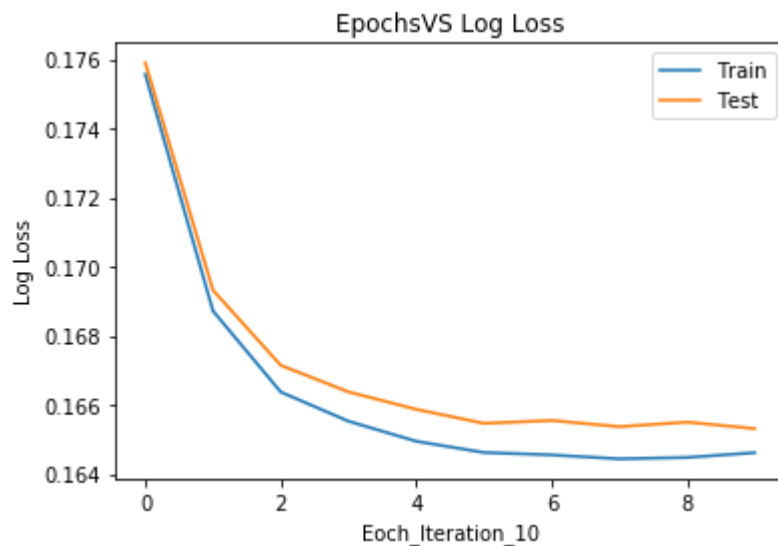
```
In [24]: # these are the results we got after we implemented sgd and found the optimal wei
w-clf.coef_, b-clf.intercept_
```

```
Out[24]: (array([[ -0.00940764,  0.00399639,  0.00877248, -0.00595807, -0.00720052,
                  0.00329988,  0.0074154 ,  0.00571249,  0.02513098,  0.00123466,
                  -0.00739057, -0.00542827,  0.00077331,  0.00118758,  0.00451163]]),
          array([0.00231397]))
```

Plot epoch number vs train , test loss

- epoch number on X-axis
- loss on Y-axis

```
In [25]: plt.plot(range(epochs), train_loss, label="Train")
plt.plot(range(epochs), test_loss, label="Test")
plt.title('EpochsVS Log Loss')
plt.xlabel('Epoch Iteration_10')
plt.ylabel('Log Loss')
plt.legend()
plt.show()
```



```
In [26]: def pred(w,b, X):
    N = len(X)
    predict = []
    for i in range(N):
        z=np.dot(w,X[i])+b
        if sigmoid(z) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
            predict.append(1)
        else:
            predict.append(0)
    return np.array(predict)
print(1-np.sum(y_train - pred(w,b,X_train))/len(X_train))
print(1-np.sum(y_test - pred(w,b,X_test))/len(X_test))
```

0.96032

0.9576

In []:

In []: